

# YASMIN

*(Yet Another Shared Memory for Intra-Node framework)*

## Efficient Intra-Node Communication Using Generic Sockets


Michalis Rozis, Stefanos Gerangelos, and Nectarios Koziris  
National Technical University of Athens  
Computing Systems, Laboratory




# What's the problem?

- VMs exchange large amount of data. For example:
  - MapReduce workloads running in many VMs.
  - HPC applications
  - Network Function Virtualisation (NFV) such as routing, encryption, load balancing etc.
  - VMs access shared device (accelerators, crypto devices etc)
- VMs can be placed in different nodes
  - Data path through network
- VMs can be placed in the same node
  - Data path through memory
  - Memory speed > Network


# Example: NFV




# Example: Shared GPU




# Example: Web Server




# Example: Co-located NFV VMs



# Example: Co-located VMs - shared GPU



# Example: Co-located VMs - Web Server





# Efficient Intra-node Communication

- Place VMs in the same node
- Make communication efficient (latency, throughput).
- Applications need to be unaware of execution environment.

# Build on the Xen hypervisor

- Widely used
- Lightweight type-1
- Paravirtualized Guests
- Exploit page sharing techniques

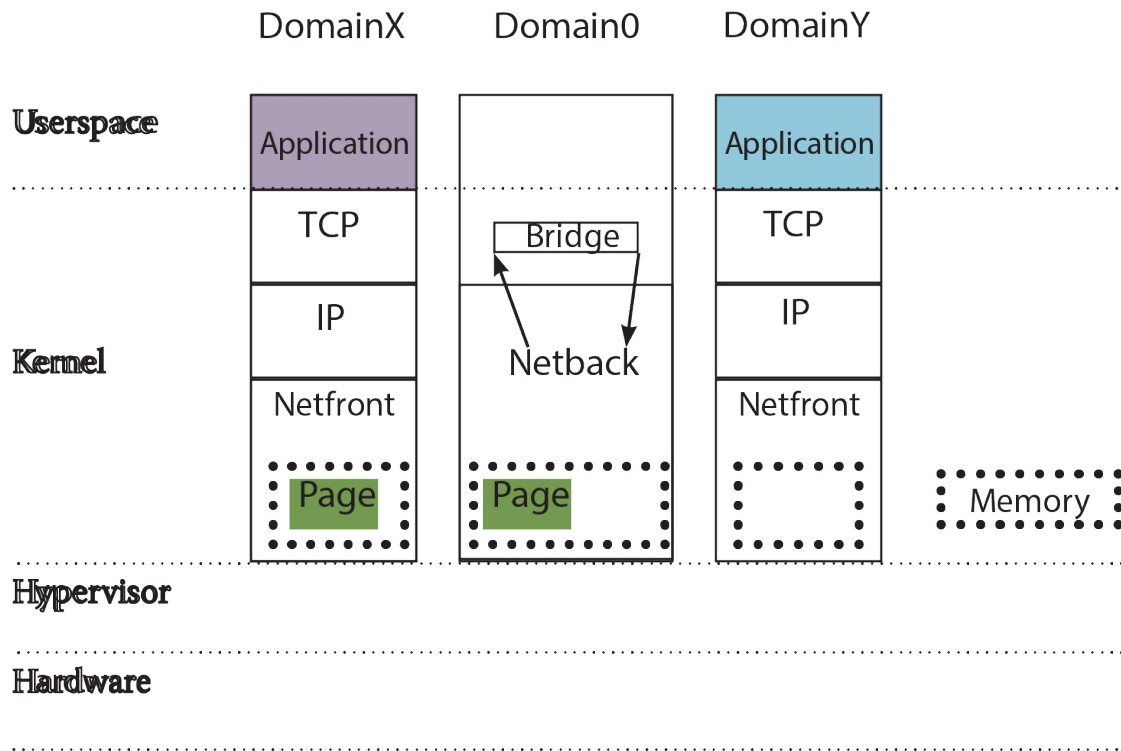
# Xen features

- Grant table → map and share pages
  - Granter (Dom1):
    - Allocates page
    - Grants page access to foreign domain (Dom2)
    - Returns grant-table index
  - Grantee (Dom2):
    - Allocates page
    - Maps page using grant-table index

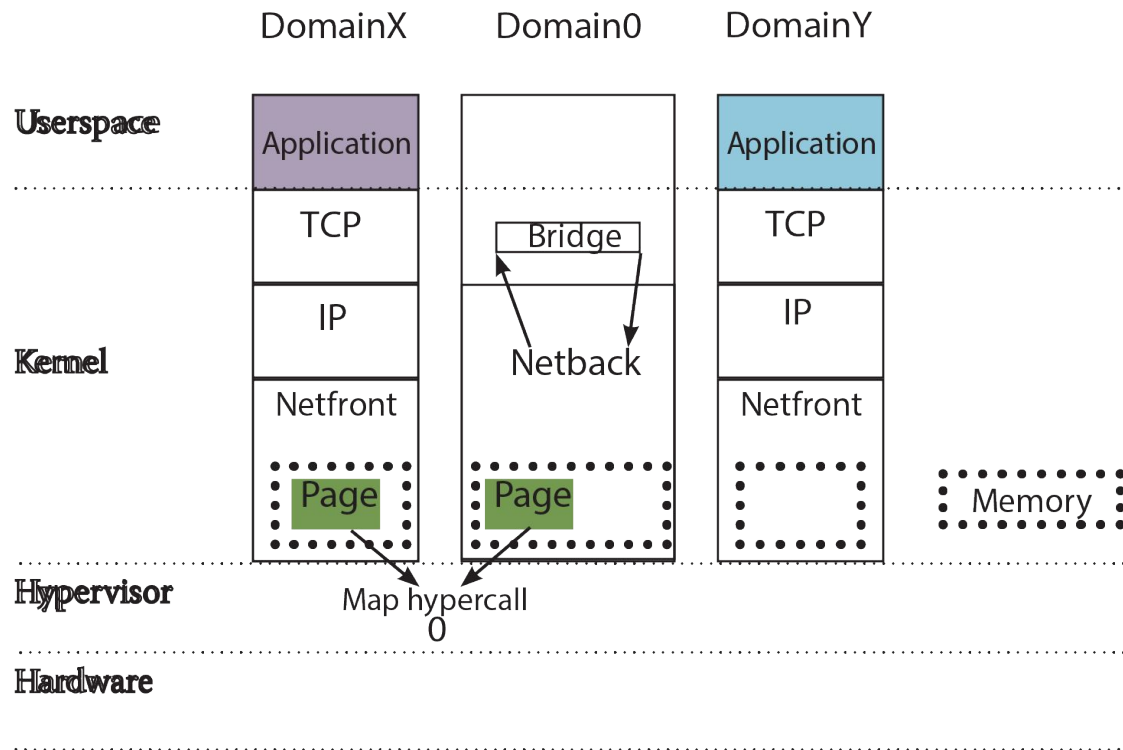
# Xen features

- Event channel → Virtual interrupt mechanism
  - Dom1:
    - New channel with foreign domain
    - Returns local channel port
    - Registers interrupt handler
  - Dom2:
    - Registers channel using Dom1's local channel port
    - Registers interrupt handler
  - Invoke interrupts between VMs.

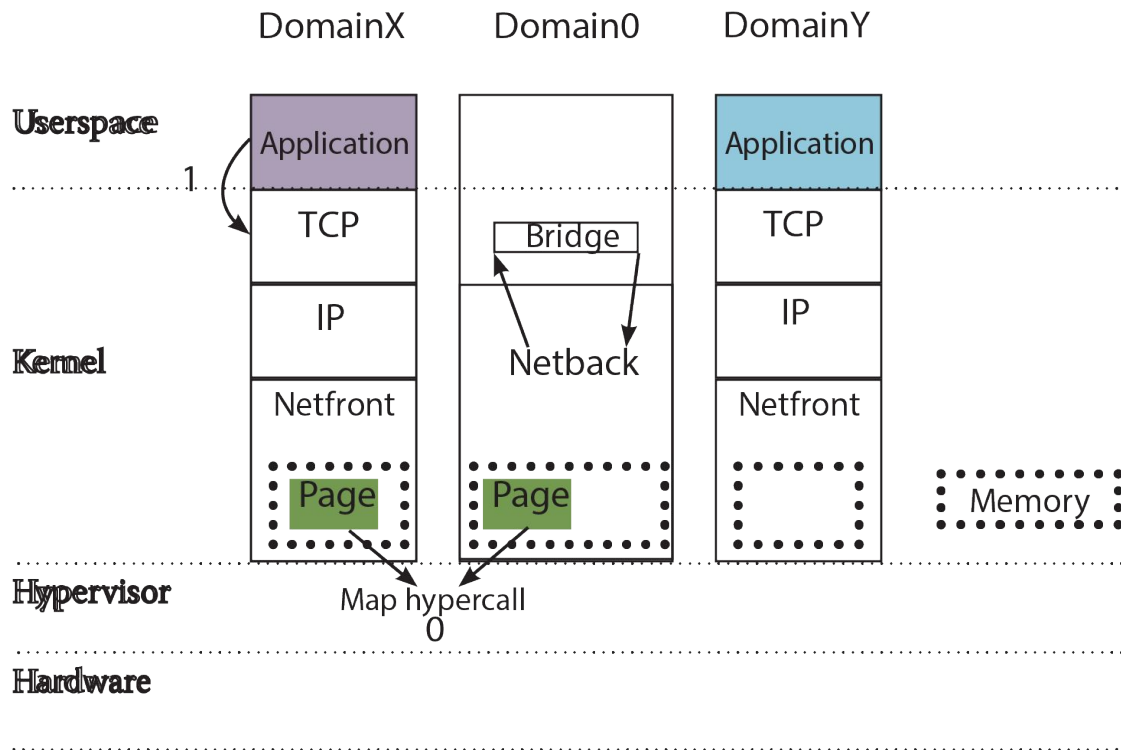
# Xen netback / netfront



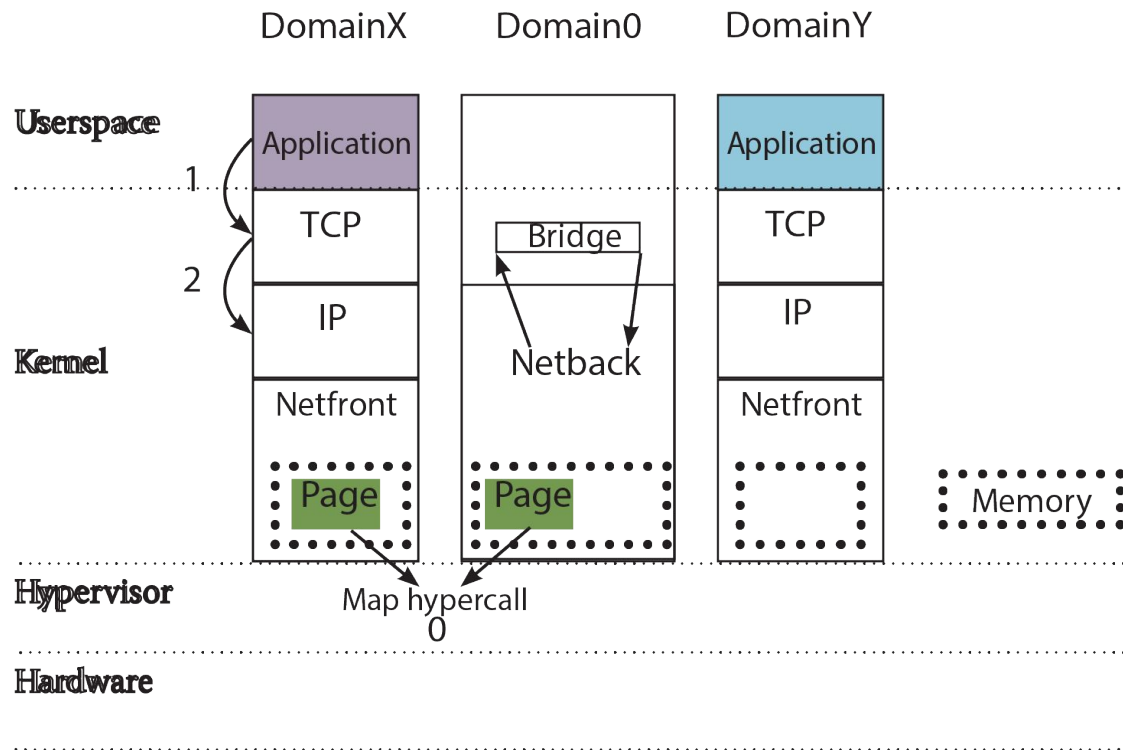
# Xen netback / netfront



# Xen netback / netfront

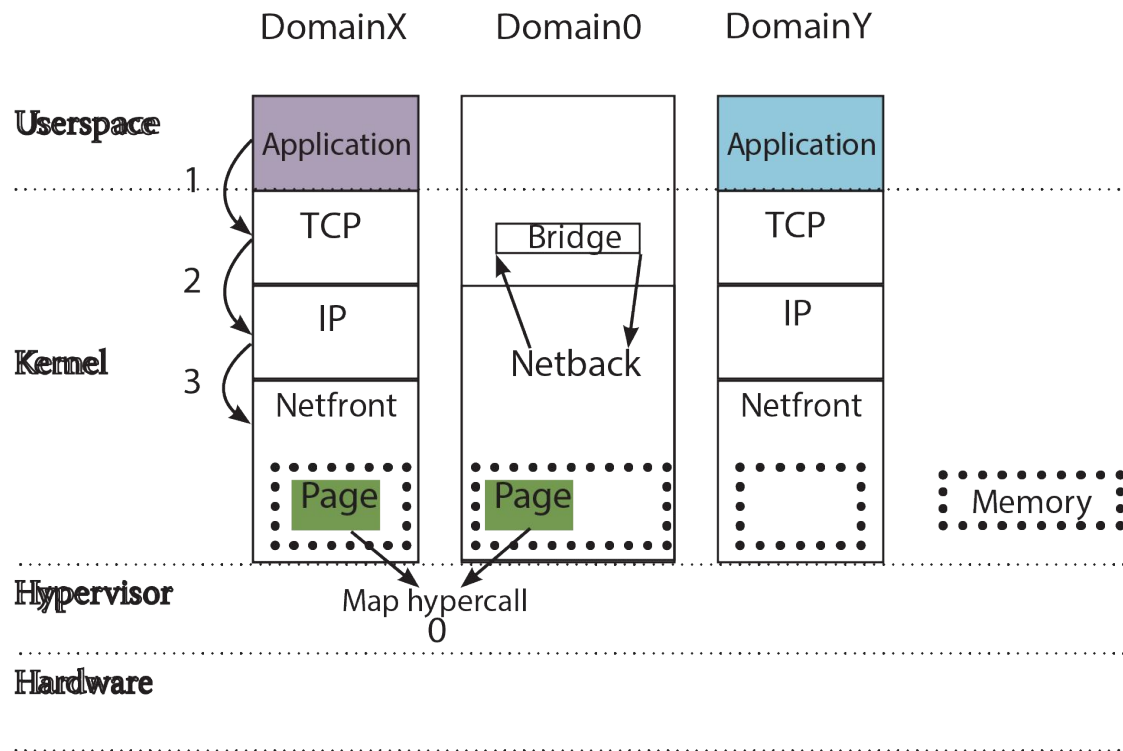


# Xen netback / netfront

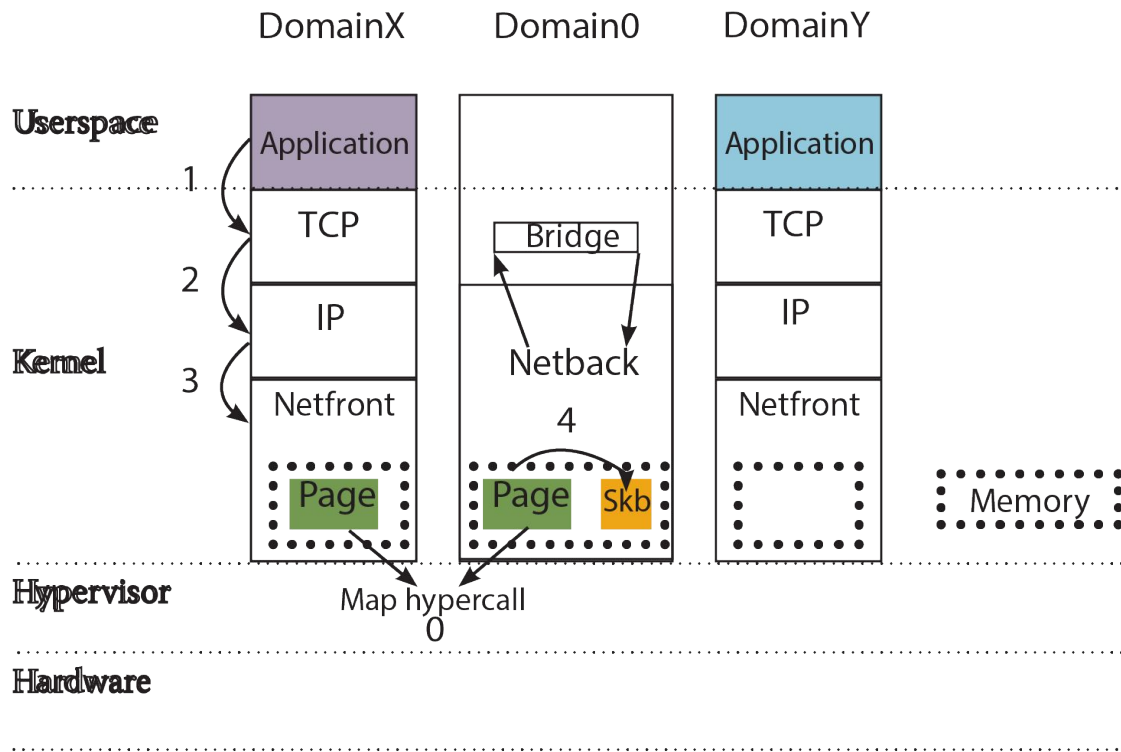




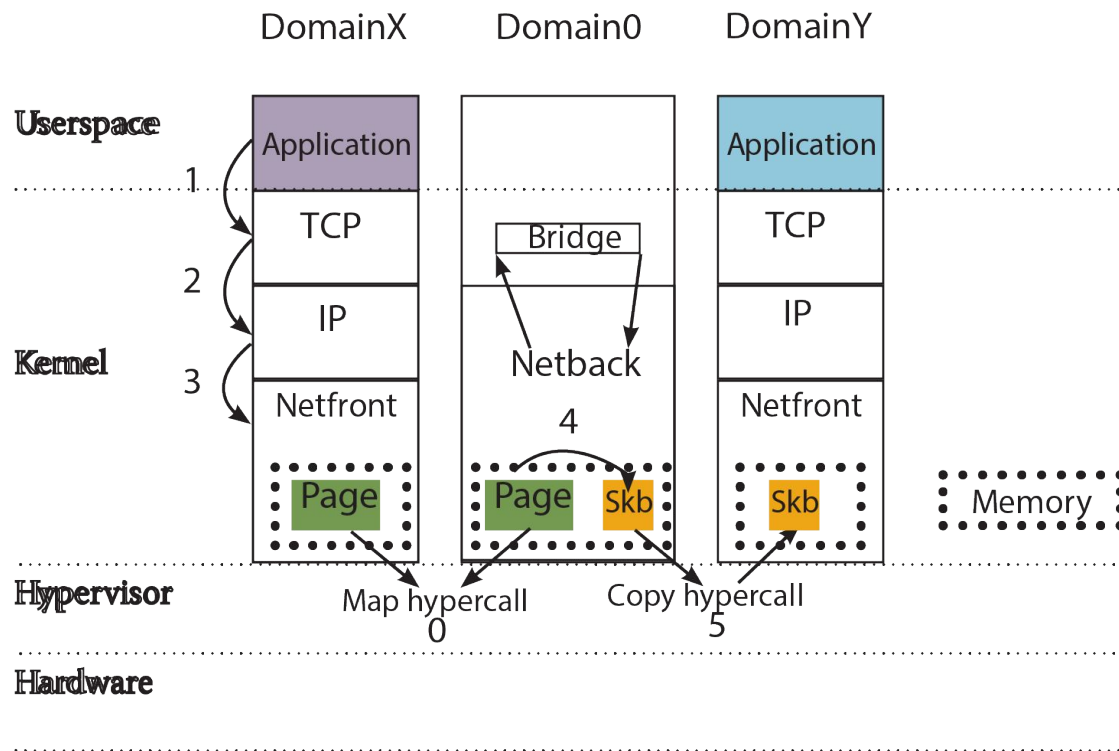
# Xen netback / netfront



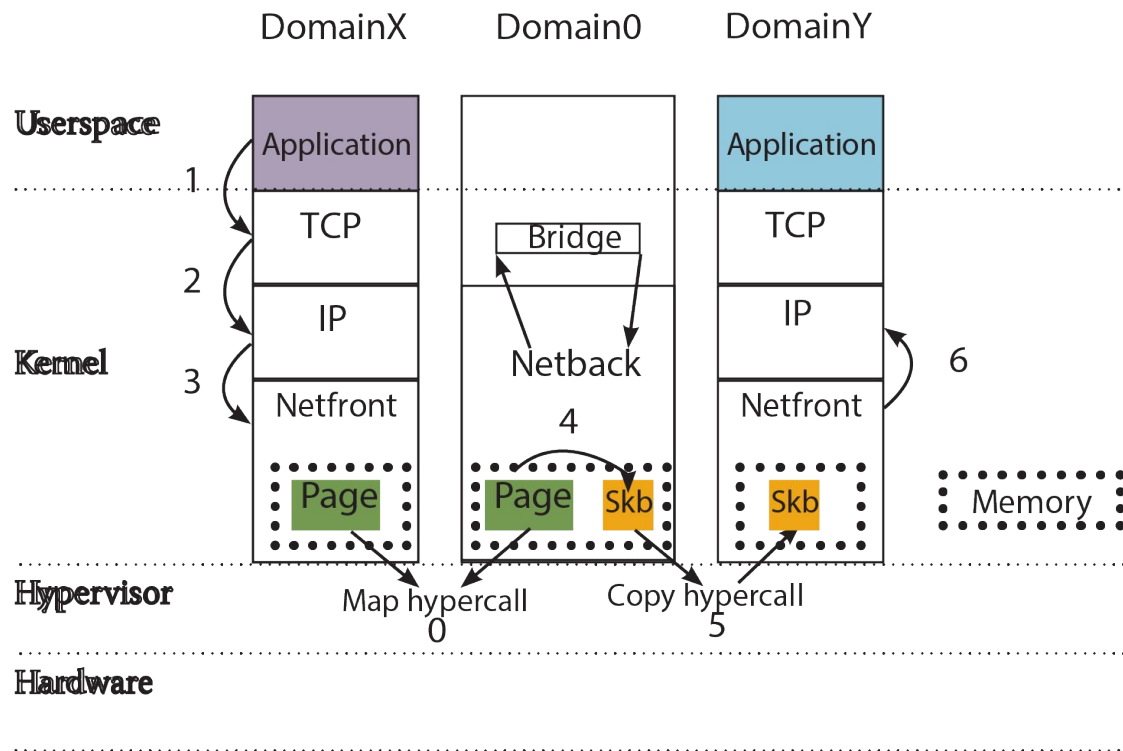
# Xen netback / netfront



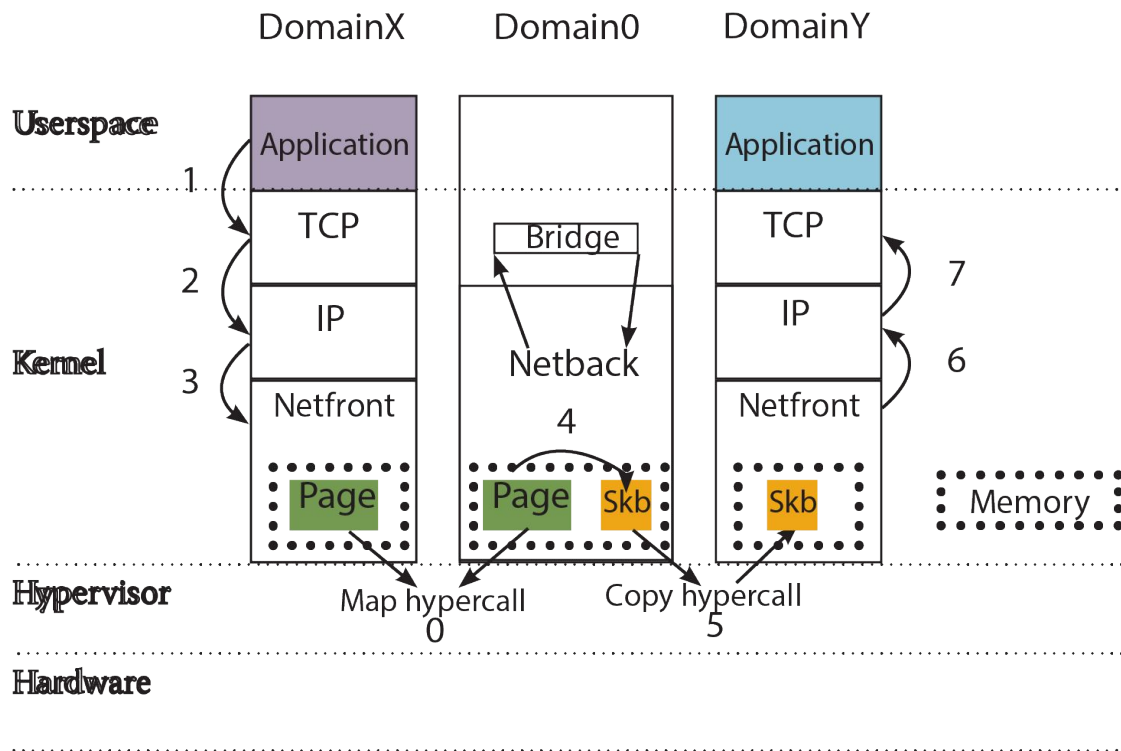
# Xen netback / netfront



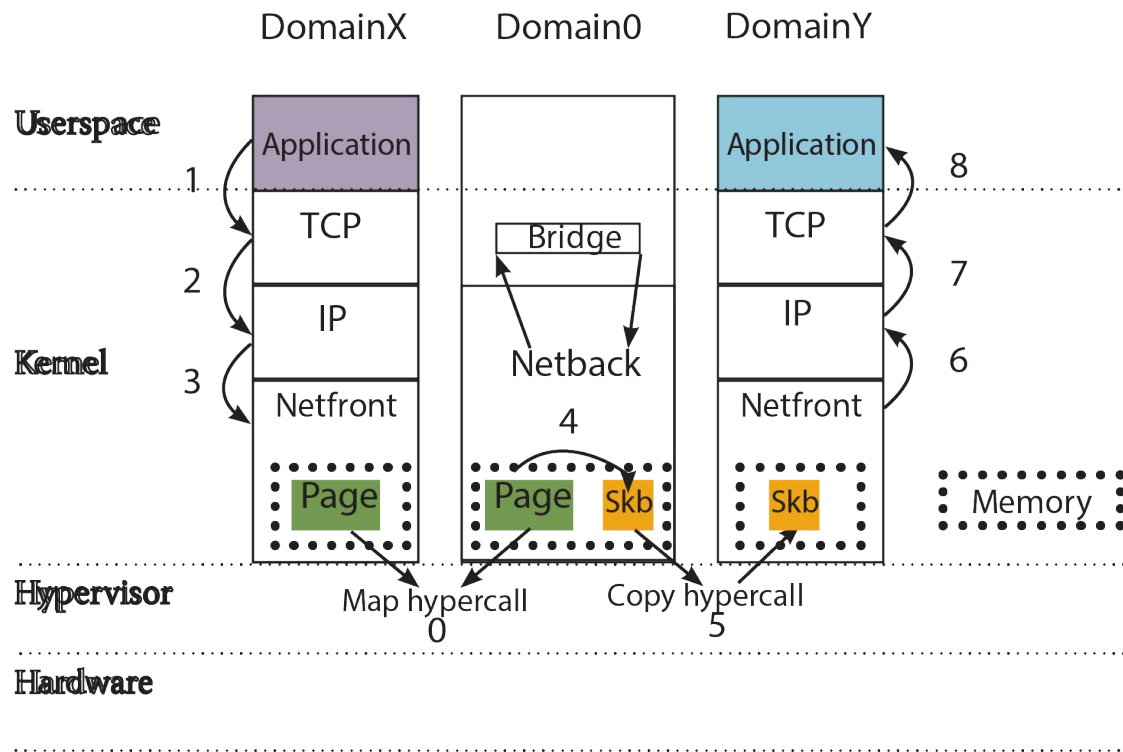
# Xen netback / netfront



# Xen netback / netfront



# Xen netback / netfront



# Standard intra-node communication:

- Slow: Copies between VMs and control domain
- Non - scalable:
  - Different applications between one pair: same interface
  - Different pair of VMs: all traffic through control domain.

✓ Leave Control Domain alone.

# How?

Idea:



- Page Sharing between connecting VMs
- Endpoints responsible for setup and breakdown of channel.
- Sockets
- Packet send/receive through network → copies to/from memory



# Related work

- Concept of page sharing with Sockets connections
  - New Socket API
  - Message copying by hypervisor
  - Packet capture and process (netfilter)

Plenty of proposals use page sharing but...

- Hypervisor modification
- Code refactoring
- Recompilation
- TCP/IP


Compatibility issues!

# What's new with YASMIN?

- Compatibility very important
  - Wiseman said just stick to POSIX Sockets
- Why always pay TCP/IP fee? Do we really need this? Take a shortcut!
- Build a new protocol from scratch?
  - Use vSockets (AF\_VSOCK) with a brand new transport layer
  - VMWare vSockets protocol (domain ID & connection port)


# Design overview

Default

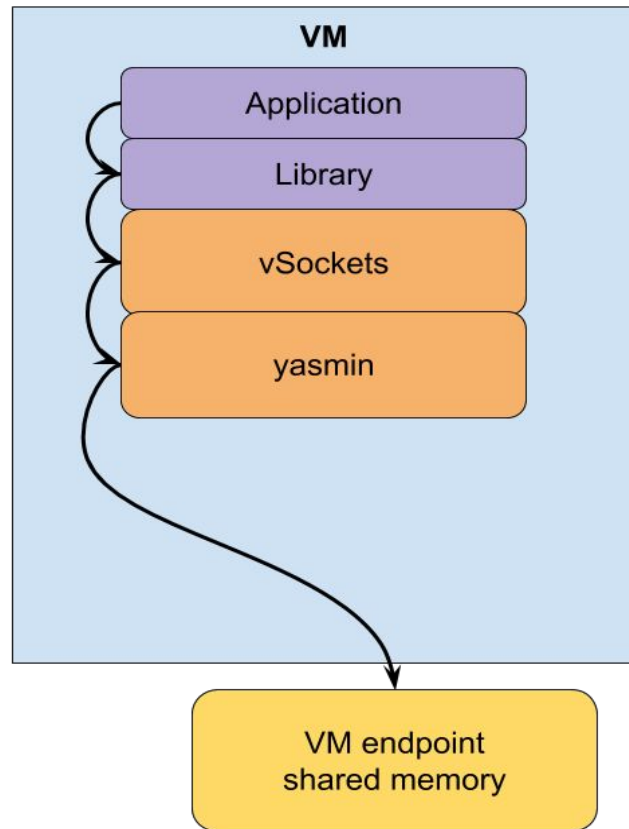


# Design overview

Default



YASMIN

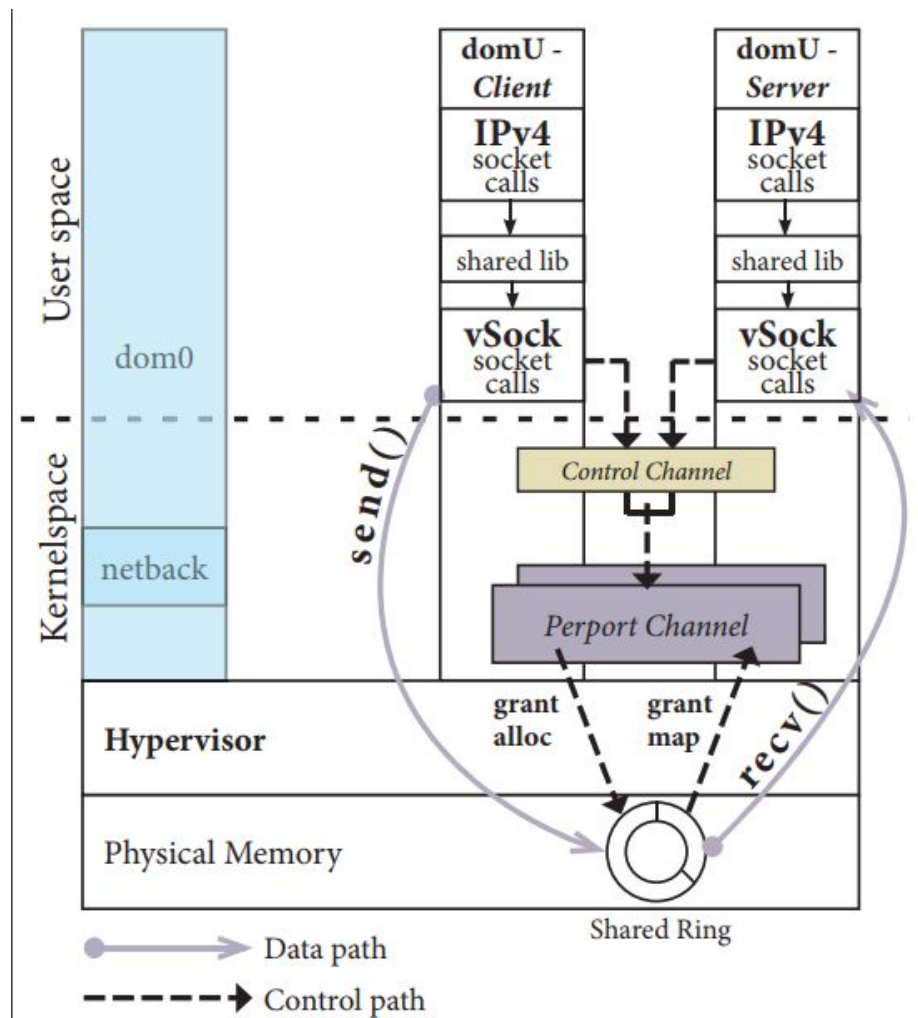


# Details

- TCP/IP socket calls translated to vSockets calls (AF\_VSOCK)
- *hosts* file with IPv4 - domID mappings
- One control channel connecting VM pairs
- One data channel for connecting socket pairs (*persocket* channel)

# Channel setup

- Grant-table mechanism for shared pages (grant access & map)
- Event-channel for packet notification
- Producer - consumer ring (no locks) in shared memory
- `copy_to/from_user()` for actual data transfer



# Evaluation

Use microbenchmarks (Iperf, NetPIPE, STREAM)

- Latency
- Bandwidth
- Scalability




# Baseline

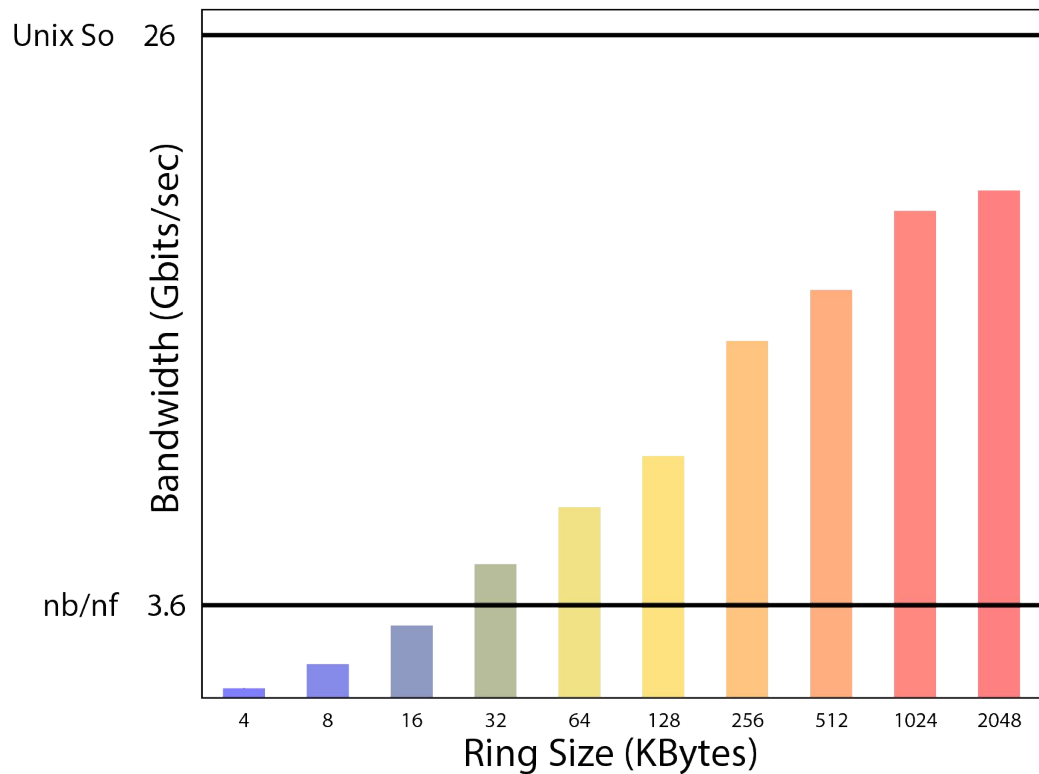
Compare to:

- netback / netfront
- Unix Domain Sockets
- Memory bandwidth

# Evaluation Setup



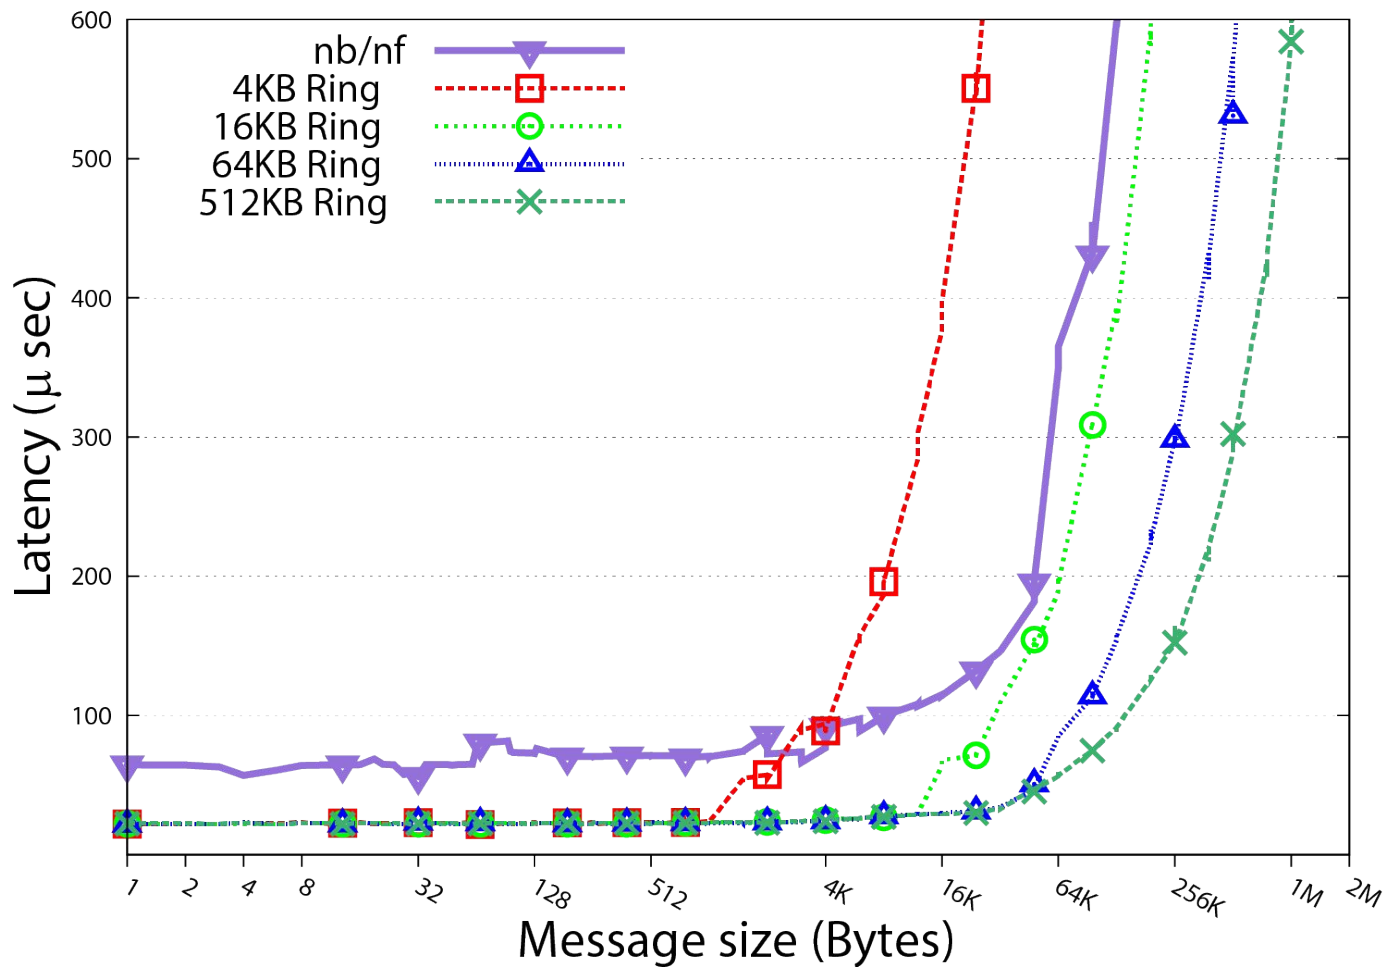
# Proper Ring Size




# Latency - throughput

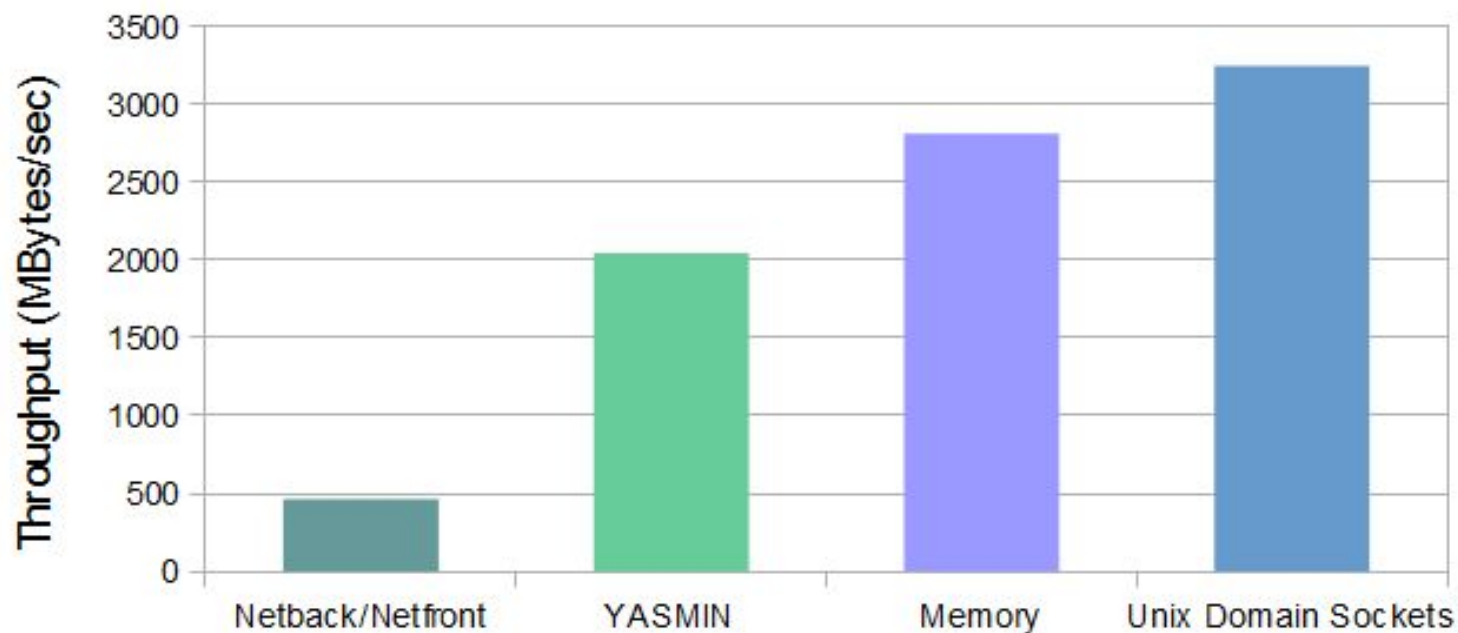
For 512kB ring size (128 pages):

- Latency reduction up to 65% (netback/netfront)
  - YASMIN @22μsec
  - netback/netfront @64μsec
- Throughput increase by 4.4x (netback/netfront)
  - netback/netfront @ 463MBytes/sec
  - YASMIN @2048MBytes/sec
  - Memory@2813MBytes/sec
  - UNIX Domain Sockets @3250MBytes/sec






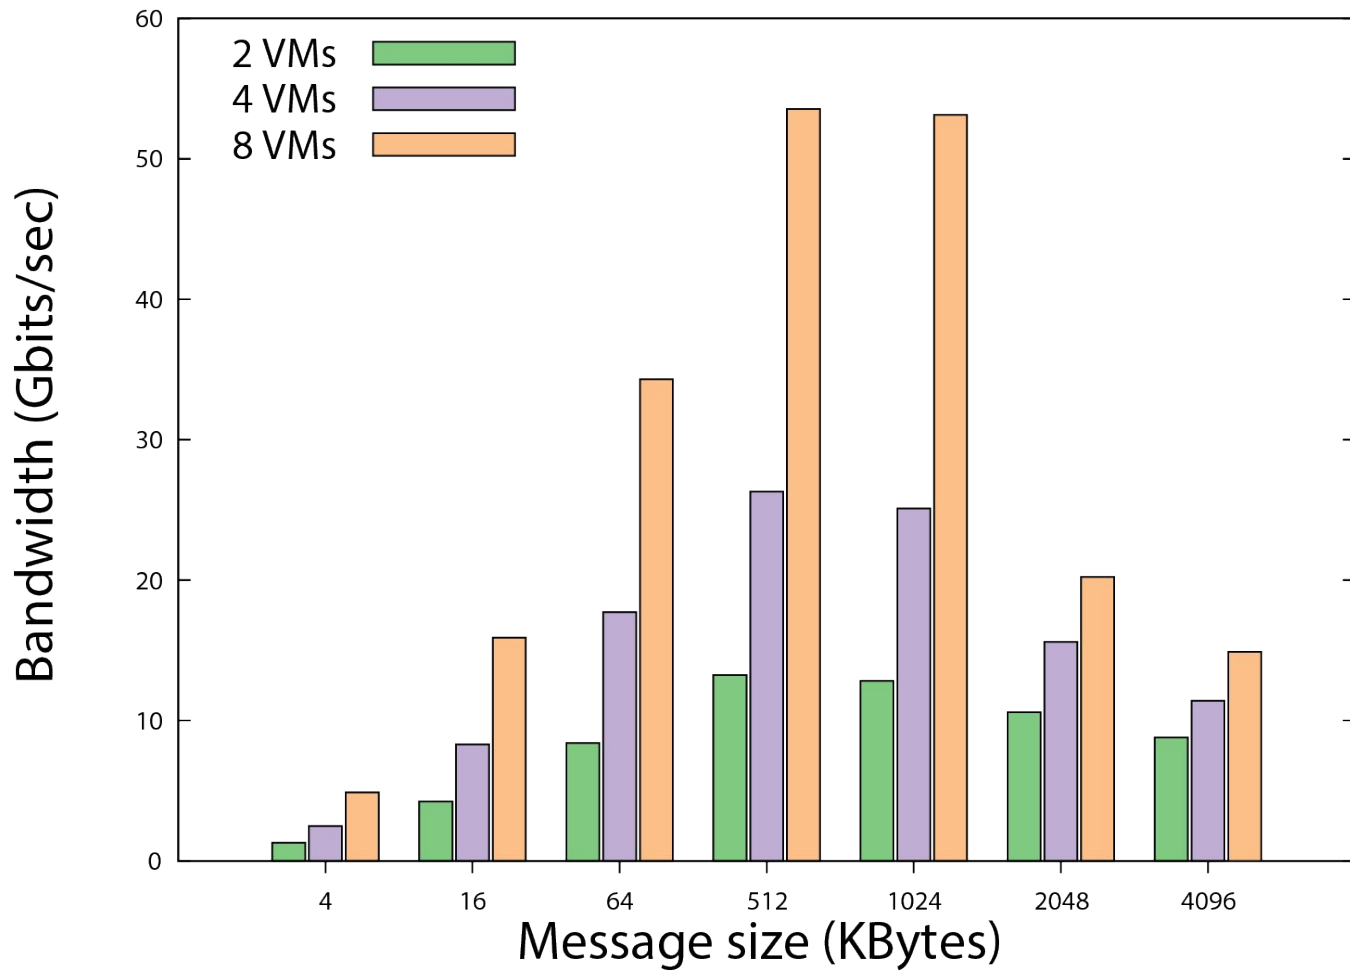
## Performance Comparison



# Scalability setup:







# Conclusion

- VMs exchange data → Run in same physical node
- POSIX Sockets - Retain compatibility
  - No refactoring
  - No recompiling
  - No hypervisor modification
- Choose transport layer on the fly
- Exploit shared memory
- Bypass TCP/IP
- Latency - Throughput - Scalability

# Future work

- Test on real-life applications (NFV infrastructure)
- Validate caching effects
- Enhance hosts file mechanism